

# Data wrangling

Using dplyr to transform your data

**Statistical Computing and Empirical Methods**  
**Unit EMATM0061, Data Science MSc**

Rihuan Ke

rihuan.ke@bristol.ac.uk

Teaching Block 1, 2024

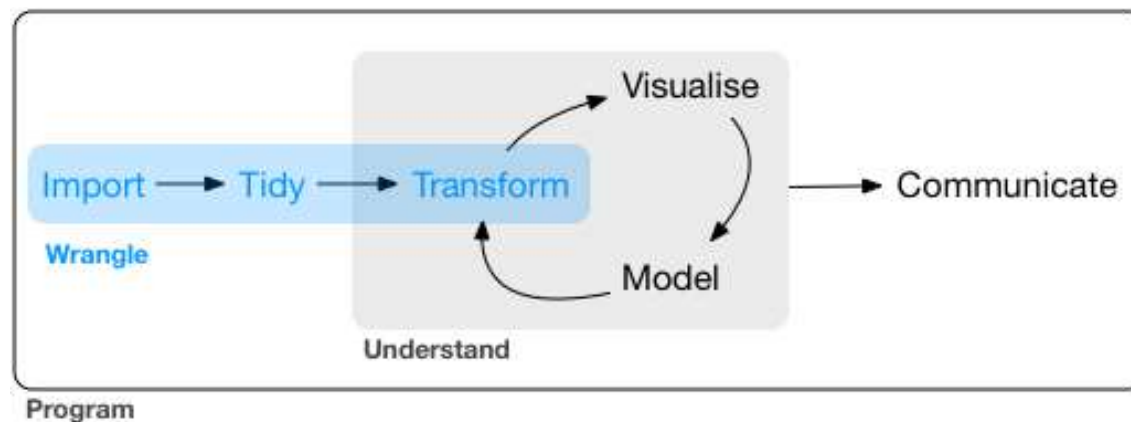
# *What we will cover in this lecture*

- We will introduce concepts in basic **data wrangling operations**
  - Select, filter, mutate, arrange, summarize, ...
- We will learn how to perform data wrangling operations using tools in the programming language R
  - The package **dplyr**

# What is data wrangling

*Data wrangling*: the process of transforming data from one form to another in preparation for another downstream task (e.g., visualisation, modelling)

Transforming your data with basic **data wrangling operations**: selecting, filtering, mutating, arranging, summarizing, joining...



source: [r4ds.had.co.nz](http://r4ds.had.co.nz)

# *Learning data wrangling with examples in R*

We will learn data wrangling with examples in R

We will do the examples with two important R packages

## *1. The dplyr package*

- An R package designed for data wrangling
- Effective data wrangling APIs, such as
  - *select()*, *filter()*, *mutate()*, ...

## *2. The tidyverse package*

- A collection of R packages that are designed for data science
- Including
  - *ggplot2*: a package for visualisation
  - *tidyr*: a package for tidying data
  - *dplyr*
  - *purrr*: functional programming

Install and load the packages:

```
install.packages("tidyverse")  
library(tidyverse)
```

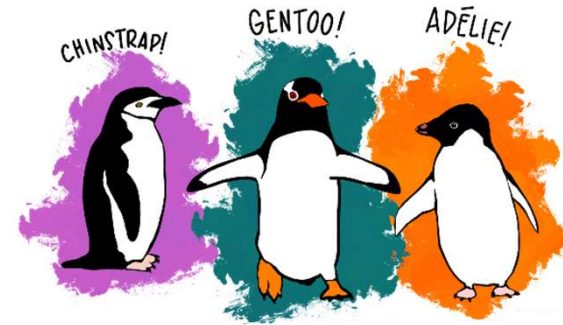
The dplyr package is included in the tidyverse package (so it is loaded when tidyverse is loaded)

# Case study: the Palmer penguins data set

Examples will be demonstrated using the *Palmer penguins data set*, which was introduced by Alison Hill, Allison Horst, Kristen Gorman

To use the Palmer penguins data set:

```
install.packages("palmerpenguins")  
library(palmerpenguins)
```



This is what the *penguins* data looks like:

```
head(penguins)
```

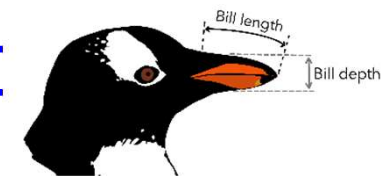
```
## # A tibble: 6 x 8  
##   species island bill_length_mm bill_depth_mm flipper_l...1 body_...2 sex   year  
##   <fct>   <fct>         <dbl>         <dbl>         <int>   <int> <fct> <int>  
## 1 Adelia Torgersen      39.1           18.7           181    3750 male  2007  
## 2 Adelia Torgersen      39.5           17.4           186    3800 fema... 2007  
## 3 Adelia Torgersen      40.3            18           195    3250 fema... 2007  
## 4 Adelia Torgersen      NA              NA              NA      NA <NA>  2007  
## 5 Adelia Torgersen      36.7           19.3           193    3450 fema... 2007  
## 6 Adelia Torgersen      39.3           20.6           190    3650 male  2007  
## # ... with abbreviated variable names 1flipper_length_mm, 2body_mass_g
```

In R, data sets are often stored as data frames

# Tabular data

*Penguins* is an example of a tabular data set represented by an R data frame.

```
## # A tibble: 6 x 8
##   species island  bill_length_mm bill_depth_mm flipper_l... body_m...² sex   year
##   <fct>   <fct>      <dbl>         <dbl>      <int>      <int> <fct> <int>
## 1 Adelie  Torgersen    39.1          18.7        181       3750 male  2007
## 2 Adelie  Torgersen    39.5          17.4        186       3800 fema... 2007
## 3 Adelie  Torgersen    40.3          18          195       3250 fema... 2007
## 4 Adelie  Torgersen    NA            NA           NA         NA <NA>  2007
## 5 Adelie  Torgersen    36.7          19.3        198       3450 fema... 2007
## 6 Adelie  Torgersen    39.3          20.6        190       3650 male  2007
## # ... with abbreviated variable names ^flipper_length_mm, ^body_mass_g
```



*rows*

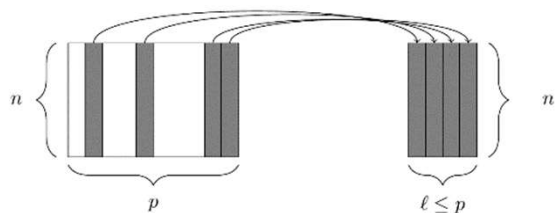
Each row corresponds to an instance of a specific type of thing, in this case, an individual penguin. Known as examples, observations or cases.

*columns*

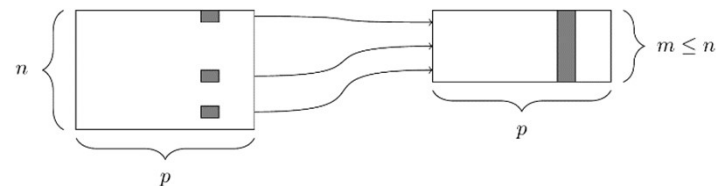
Each column (also called a **variable**) corresponds to a property or quality of the individual examples. Known as features, or variables.

In the Penguins data set, we have 8 columns, corresponding to different properties of a penguin: “*species*”, “*island*”, “*bill\_length\_mm*”, “*bill\_depth\_mm*”, “*flipper\_length\_mm*”, “*body\_mass\_g*”, “*sex*”, “*year*”

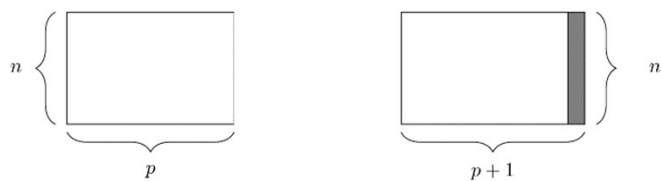
# Data wrangling operations



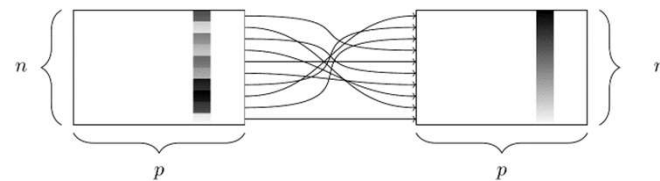
select



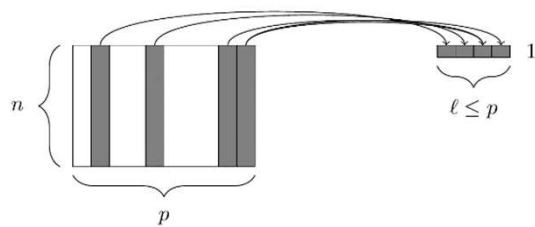
filter



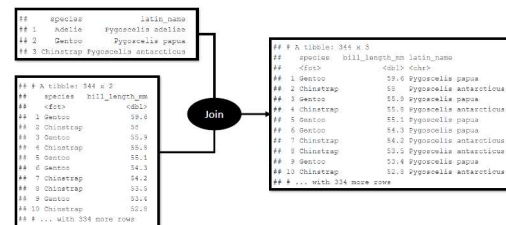
Create new columns



sort

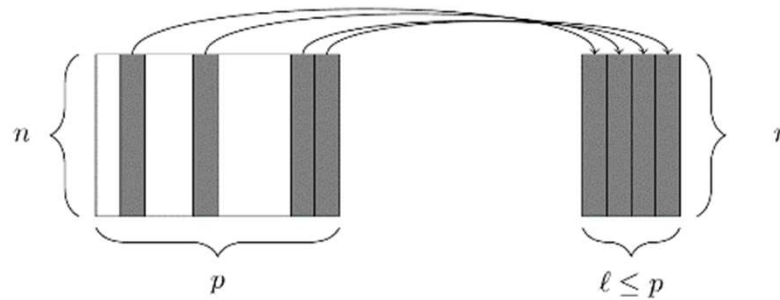


summarize



join

# 1. Selecting columns



Selecting a subset of columns and generating a new dataset (with fewer columns)

In R, this can be done with the `select()` function (from the dplyr package), e.g.,

```
penguinsv2 <- select(penguins, species, bill_length_mm, body_mass_g, flipper_length_mm )
print(penguinsv2)
```

```
## # A tibble: 344 x 4
##   species bill_length_mm body_mass_g flipper_length_mm
##   <fct>      <dbl>      <int>      <int>
## 1 Adelie      39.1      3750      181
## 2 Adelie      39.5      3800      186
## 3 Adelie      40.3      3250      195
## 4 Adelie      NA         NA         NA
## 5 Adelie      36.7      3450      193
## 6 Adelie      39.3      3650      190
## 7 Adelie      38.9      3625      181
## 8 Adelie      39.2      4675      195
## 9 Adelie      34.1      3475      193
## 10 Adelie     42       4250      190
## # ... with 334 more rows
```

The result `penguinsv2` is a new data frame (with 4 columns), which we will use frequently in the following examples



# 1. Selecting columns

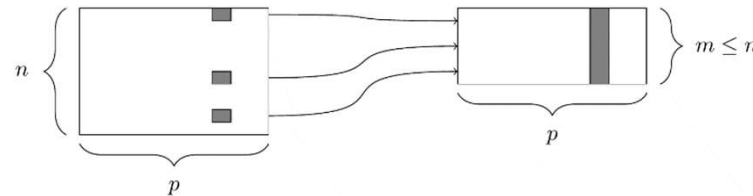
The select function also allows us to **remove several columns** (invert selection) using the symbol '-', e.g.,

```
select(penguins, -species, -bill_length_mm, -body_mass_g)
```

```
## # A tibble: 344 x 5
##   island    bill_depth_mm flipper_length_mm sex    year
##   <fct>         <dbl>             <int> <fct> <int>
## 1 Torgersen      18.7               181 male   2007
## 2 Torgersen      17.4               186 female 2007
## 3 Torgersen      18                195 female 2007
## 4 Torgersen      NA                NA <NA>   2007
## 5 Torgersen      19.3               193 female 2007
## 6 Torgersen      20.6               190 male   2007
## 7 Torgersen      17.8               181 female 2007
## 8 Torgersen      19.6               195 male   2007
## 9 Torgersen      18.1               193 <NA>   2007
## 10 Torgersen     20.2               190 <NA>   2007
## # ... with 334 more rows
```

So we have 5 columns (after removing the three columns)

## 2. Filtering rows

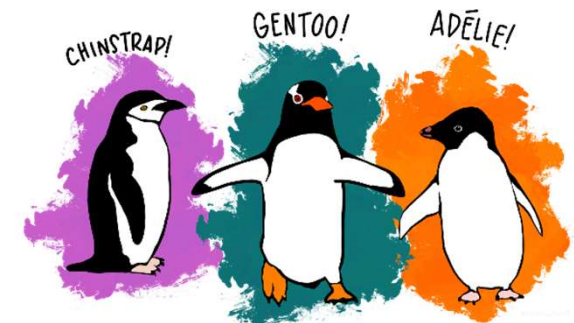


Extracting a subset of rows (while the columns are unchanged)

In R, this can be done with the `filter()` function (from the dplyr package), e.g.,

```
filter(penguinsv2, species=='Gentoo')
```

```
## # A tibble: 124 x 4
##   species bill_length_mm body_mass_g flipper_length_mm
##   <fct>         <dbl>         <int>         <int>
## 1 Gentoo         46.1           4500           211
## 2 Gentoo          50           5700           230
## 3 Gentoo         48.7           4450           210
## 4 Gentoo          50           5700           218
## 5 Gentoo         47.6           5400           215
## 6 Gentoo         46.5           4550           210
## 7 Gentoo         45.4           4800           211
## 8 Gentoo         46.7           5200           219
## 9 Gentoo         43.3           4400           209
## 10 Gentoo        46.8           5150           215
## # ... with 114 more rows
```



So we get rows associated with penguins that are of 'Gentoo' species

## 2. Filtering rows

We can select the rows that satisfy **multiple conditions** (using the expression '&')

For example, to select penguins that are of the Gentoo species and has body mass bigger than 5kg:

```
filter(penguinsv2, species=='Gentoo' & body_mass_g>5000)
```

```
## # A tibble: 61 x 4
##   species bill_length_mm body_mass_g flipper_length_mm
##   <fct>         <dbl>         <int>         <int>
## 1 Gentoo         50           5700           230
## 2 Gentoo         50           5700           218
## 3 Gentoo        47.6           5400           215
## 4 Gentoo        46.7           5200           219
## 5 Gentoo        46.8           5150           215
## 6 Gentoo        49           5550           216
## 7 Gentoo        48.4           5850           213
## 8 Gentoo        49.3           5850           217
## 9 Gentoo        49.2           6300           221
## 10 Gentoo       48.7           5350           222
## # ... with 51 more rows
```

# Combining filter & select functions

The functions *select* and *filter* can be used together (to select a subset of columns and a subset of rows)

```
select(filter(penguinsv2, species=='Gentoo'), species, bill_length_mm, body_mass_g)
```

```
## # A tibble: 124 × 3
##   species bill_length_mm body_mass_g
##   <fct>         <dbl>         <int>
## 1 Gentoo         46.1           4500
## 2 Gentoo          50           5700
## 3 Gentoo         48.7           4450
## 4 Gentoo          50           5700
## 5 Gentoo         47.6           5400
## 6 Gentoo         46.5           4550
## 7 Gentoo         45.4           4800
## 8 Gentoo         46.7           5200
## 9 Gentoo         43.3           4400
## 10 Gentoo        46.8           5150
## # ... with 114 more rows
```

So we get only three columns & rows that are associated with “Gentoo” species

# *Simplifying codes with the pipe operator*

We can also chain multiple operations with the **pipe operator %>%**

The following statements are equivalent:

```
select(filter(penguinsv2, species=='Gentoo'), species, bill_length_mm, body_mass_g)
```

```
penguinsv2 %>%  
  filter(species=='Gentoo') %>%  
  select(species, bill_length_mm, body_mass_g)
```

The pipe operator %>% allows arguments to be implicitly passed as objects to the function after the pipe.

```
f <- function(a,b) {return (a^2 + b) }  
print(f(3,1))
```

```
## [1] 10
```

```
print( 3 %>% f(1) )
```

```
## [1] 10
```

# *Simplifying codes with the pipe operator*

To chain multiple operations (e.g., f1, f2, f3), we have

`x %>% f1(a) %>% f2(b) %>% f3(c)` means `f3(f2(f1(x, a), b), c)`

The pipe operator `%>%` is taken from the *magrittr* package which is also part of the *tidyverse*

The *magrittr* package was developed by Stefan Milton Bache and Hadley Wickham.

# 3. Creating and renaming columns

The aim is to **create a new column** as a function of existing columns.



In R, this can be done with the **mutate()** function (from the dplyr package), e.g.,

```
penguinsv2 %>%  
  mutate(flipper_bill_ratio=flipper_length_mm/bill_length_mm)
```

```
## # A tibble: 344 x 5  
##   species bill_length_mm body_mass_g flipper_length_mm flipper_bill_ratio  
##   <fct>      <dbl>      <int>          <int>          <dbl>  
## 1 Adelie      39.1        3750            181            4.63  
## 2 Adelie      39.5        3800            186            4.71  
## 3 Adelie      40.3        3250            195            4.84  
## 4 Adelie      NA           NA              NA             NA  
## 5 Adelie      36.7        3450            193            5.26  
## 6 Adelie      39.3        3650            190            4.83  
## 7 Adelie      38.9        3625            181            4.65  
## 8 Adelie      39.2        4675            195            4.97  
## 9 Adelie      34.1        3475            193            5.66  
## 10 Adelie     42          4250            190            4.52  
## # ... with 334 more rows
```

A new column `flipper_bill_ratio` has been created

### 3. *Creating and renaming columns*

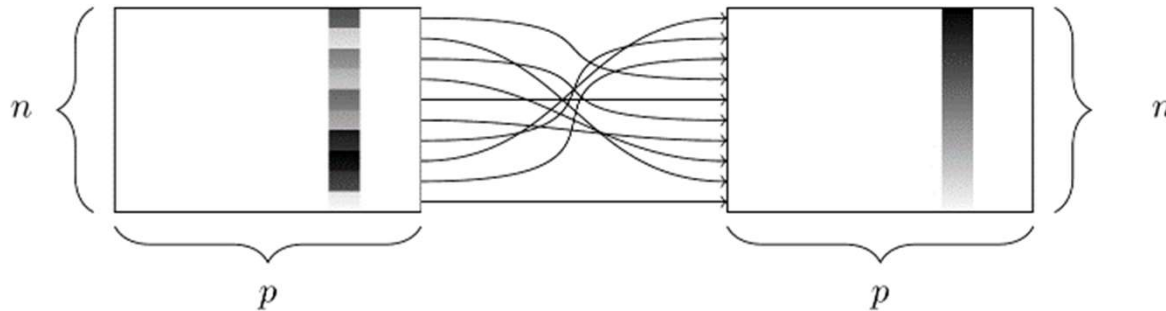
To **rename an existing column**, we can use the `rename()` function, e.g.,

```
penguinsv2 %>% rename(f_l_m = flipper_length_mm)
```

```
## # A tibble: 344 x 4
##   species bill_length_mm body_mass_g f_l_m
##   <fct>      <dbl>      <int> <int>
## 1 Adelie      39.1        3750   181
## 2 Adelie      39.5        3800   186
## 3 Adelie      40.3        3250   195
## 4 Adelie      NA           NA     NA
## 5 Adelie      36.7        3450   193
## 6 Adelie      39.3        3650   190
## 7 Adelie      38.9        3625   181
## 8 Adelie      39.2        4675   195
## 9 Adelie      34.1        3475   193
## 10 Adelie      42         4250   190
## # ... with 334 more rows
```



## 4. *Sorting the rows*



Sorting the rows of a data frame according to the values of a column

In R, this can be done with the `arrange()` function, e.g.,

```
penguinsv2 %>% arrange(bill_length_mm)
```

```
## # A tibble: 344 x 4
##   species bill_length_mm body_mass_g flipper_length_mm
##   <fct>      <dbl>      <int>      <int>
## 1 Adelie      32.1        3050         188
## 2 Adelie      33.1        2900         178
## 3 Adelie      33.5        3600         190
## 4 Adelie      34         3400         185
## 5 Adelie      34.1        3475         193
## 6 Adelie      34.4        3325         184
## 7 Adelie      34.5        2900         187
## 8 Adelie      34.6        4400         198
## 9 Adelie      34.6        3200         189
## 10 Adelie     35         3450         190
## # ... with 334 more rows
```

Now the rows are sorted according to `bill_length_mm` (in ascending order)

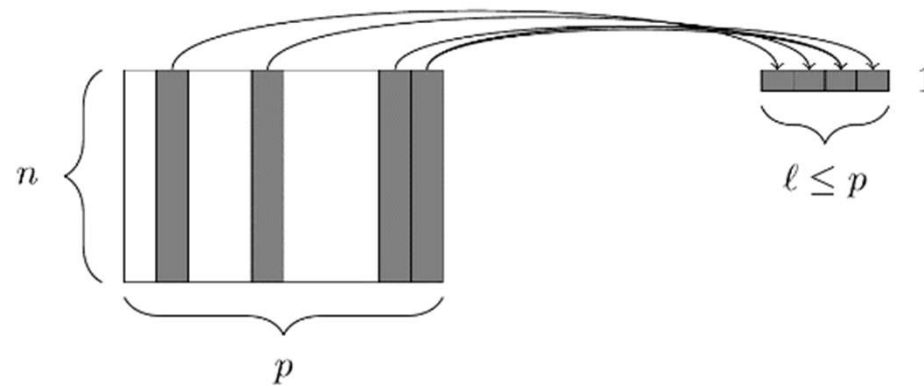
## 4. *Sorting the rows*

We can also sort in **descending order**, with `desc()`

```
penguinsv2 %>% arrange(desc(bill_length_mm))
```

```
## # A tibble: 344 x 4
##   species  bill_length_mm body_mass_g flipper_length_mm
##   <fct>         <dbl>         <int>         <int>
## 1 Gentoo         59.6           6050           230
## 2 Chinstrap      58             3700           181
## 3 Gentoo         55.9           5600           228
## 4 Chinstrap      55.8           4000           207
## 5 Gentoo         55.1           5850           230
## 6 Gentoo         54.3           5650           231
## 7 Chinstrap      54.2           4300           201
## 8 Chinstrap      53.5           4500           205
## 9 Gentoo         53.4           5500           219
## 10 Chinstrap     52.8           4550           205
## # ... with 334 more rows
```

## 5. Summarizing data



**Summarizing a data frame** into just one value or a vector (e.g., compute the mean, median, sum, standard deviation, ... of a column)

In R, this can be done with the **summarize()** function, e.g.,

```
penguinsv2 %>%  
  summarize(num_rows=n(), avg_weight_kg=mean(body_mass_g/1000, na.rm=TRUE), avg_flipper_bill_ratio=mean(flipper_length_mm/bill_length_mm, na.rm=TRUE))
```

```
## # A tibble: 1 × 3  
##   num_rows avg_weight_kg avg_flipper_bill_ratio  
##   <int>      <dbl>          <dbl>  
## 1     344        4.20            4.62
```

Here we have extracted three statistics including the *number of rows*, *average weights*, and *average flipper bill ratio* (which are contained in the three columns of the output)

# Group the rows when summarizing

We can use the function `group_by()` to group the rows of the data frames according to some given criteria, e.g. species

```
penguinsv2 %>%  
  group_by(species)
```

```
## # A tibble: 344 x 4  
## # Groups:   species [3]  
##   species bill_length_mm body_mass_g flipper_length_mm  
##   <fct>         <dbl>         <int>         <int>  
## 1 Adelie         39.1           3750           181  
## 2 Adelie         39.5           3800           186  
## 3 Adelie         40.3           3250           195  
## 4 Adelie         NA              NA             NA  
## 5 Adelie         36.7           3450           193  
## 6 Adelie         39.3           3650           190  
## 7 Adelie         38.9           3625           181  
## 8 Adelie         39.2           4675           195  
## 9 Adelie         34.1           3475           193  
## 10 Adelie        42            4250           190  
## # ... with 334 more rows
```

# Summarize by group

Group and then summarize:

```
penguinsv2 %>%  
  group_by(species) %>%  
  summarize(num_rows=n(), avg_weight_kg=mean(body_mass_g/1000, na.rm=TRUE), avg_flipper_bill_ratio=mean(flipper_length_mm/bill_length_mm, na.rm=TRUE))
```

```
## # A tibble: 3 x 4  
##   species  num_rows avg_weight_kg avg_flipper_bill_ratio  
##   <fct>      <int>      <dbl>          <dbl>  
## 1 Adelie      152        3.70           4.92  
## 2 Chinstrap   68        3.73           4.02  
## 3 Gentoo     124        5.08           4.58
```

Now we have extracted the three statistics for individual groups (instead of the whole data frame)

# Column-wise operations with across()

Suppose that we want to compute the number of NA (not available) values in each column, which can be done via

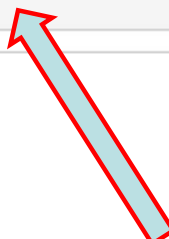
```
Num_NAs <- penguinsv2 %>% summarize(species=sum(is.na(species)), bill_length_mm=sum(is.na(bill_length_mm)), body_mass_g=sum(is.na(body_mass_g)), flipper_length_mm=sum(is.na(flipper_length_mm)))  
print(Num_NAs)
```

```
## # A tibble: 1 x 4  
##   species bill_length_mm body_mass_g flipper_length_mm  
##   <int>      <int>      <int>      <int>  
## 1      0          2          2          2
```

Use across to perform **column-wise operations** (for all columns), without copying and pasting the same code (e.g., `sum(is.na(species))`, ...)

```
Num_NAs <- penguinsv2 %>% summarize(across(everything(), ~sum(is.na(.x))))  
print(Num_NAs)
```

```
## # A tibble: 1 x 4  
##   species bill_length_mm body_mass_g flipper_length_mm  
##   <int>      <int>      <int>      <int>  
## 1      0          2          2          2
```



Here, `~sum(is.na(.x))` is equivalent to `function(x){(sum(is.na(x)))}`

# Column-wise operations on a subset of columns

We can combine `across()` and `where()` to perform **column-wise operations for a subset of columns** (for example, that is of numeric type)

```
penguinsv2 %>% summarize(across(where(is.numeric), ~mean(.x, na.rm=TRUE)))
```


```
## # A tibble: 1 × 3  
##   bill_length_mm body_mass_g flipper_length_mm  
##           <dbl>         <dbl>           <dbl>  
## 1           43.9         4202.             201.
```

# Column-wise summarizing by groups

We can combine the summarize, group\_by and across functions to perform *Column-wise summarizing by groups*

```
penguinsv2 %>%  
  select(-bill_length_mm) %>%  
  group_by(species) %>%  
  summarize(across(where(is.numeric), ~mean(.x, na.rm=TRUE)))
```

```
## # A tibble: 3 x 3  
##   species    body_mass_g flipper_length_mm  
##   <fct>         <dbl>         <dbl>  
## 1 Adelie       3701.           190.  
## 2 Chinstrap   3733.           196.  
## 3 Gentoo      5076.           217.
```



Here, **~mean(.x, na.rm=TRUE)**  
is equivalent to  
**function(x){(mean(x, na.rm=TRUE))}**

Here, we compute the mean of a column which is of numeric type, for each species of penguins



# 6. *Joining multiple data frames*

## Combining multiple data frames

Data frame 1

```
##   species      latin_name
## 1  Adelie      Pygoscelis adeliae
## 2  Gentoo      Pygoscelis papua
## 3 Chinstrap Pygoscelis antarcticus
```

```
## # A tibble: 344 x 2
##   species bill_length_mm
##   <fct>      <dbl>
## 1 Gentoo      59.6
## 2 Chinstrap    58
## 3 Gentoo      55.9
## 4 Chinstrap    55.8
## 5 Gentoo      55.1
## 6 Gentoo      54.3
## 7 Chinstrap    54.2
## 8 Chinstrap    53.5
## 9 Gentoo      53.4
## 10 Chinstrap   52.8
## # ... with 334 more rows
```

Data frame 2

Join

```
## # A tibble: 344 x 3
##   species bill_length_mm latin_name
##   <fct>      <dbl> <chr>
## 1 Gentoo      59.6 Pygoscelis papua
## 2 Chinstrap    58 Pygoscelis antarcticus
## 3 Gentoo      55.9 Pygoscelis papua
## 4 Chinstrap    55.8 Pygoscelis antarcticus
## 5 Gentoo      55.1 Pygoscelis papua
## 6 Gentoo      54.3 Pygoscelis papua
## 7 Chinstrap    54.2 Pygoscelis antarcticus
## 8 Chinstrap    53.5 Pygoscelis antarcticus
## 9 Gentoo      53.4 Pygoscelis papua
## 10 Chinstrap   52.8 Pygoscelis antarcticus
## # ... with 334 more rows
```

New data frame

In R, this can be done with **the join functions** (a part of the dplyr package)

# Example

First, create data frame 1: a data frame of bill lengths and species.

```
penguin_bill_lengths_df <- penguinsv2 %>%  
  arrange(desc(bill_length_mm)) %>%  
  select(species, bill_length_mm)  
penguin_bill_lengths_df
```

```
## # A tibble: 344 × 2  
##   species    bill_length_mm  
##   <fct>         <dbl>  
## 1 Gentoo         59.6  
## 2 Chinstrap      58  
## 3 Gentoo         55.9  
## 4 Chinstrap      55.8  
## 5 Gentoo         55.1  
## 6 Gentoo         54.3  
## 7 Chinstrap      54.2  
## 8 Chinstrap      53.5  
## 9 Gentoo         53.4  
## 10 Chinstrap     52.8  
## # ... with 334 more rows
```

Here we have used the functions `arrange` and `select` that were introduced before.

# Example

Second, create data frame 2: a data frame of Latin species names

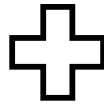
```
species <- unique(penguinsv2$species)
latin_name <- c('Pygoscelis adeliae', 'Pygoscelis papua', 'Pygoscelis antarcticus')
latin_name_df <- data.frame( species, latin_name )
print(latin_name_df)
```

```
##      species      latin_name
## 1   Adelie Pygoscelis adeliae
## 2   Gentoo Pygoscelis papua
## 3 Chinstrap Pygoscelis antarcticus
```

# Example

Data frame 1

```
## # A tibble: 344 × 2
##   species    bill_length_mm
##   <fct>         <dbl>
## 1 Gentoo         59.6
## 2 Chinstrap      58
## 3 Gentoo         55.9
## 4 Chinstrap      55.8
## 5 Gentoo         55.1
## 6 Gentoo         54.3
## 7 Chinstrap      54.2
## 8 Chinstrap      53.5
## 9 Gentoo         53.4
## 10 Chinstrap     52.8
## # ... with 334 more rows
```



Data frame 2

```
##   species    latin_name
## 1  Adelie    Pygoscelis adeliae
## 2   Gentoo    Pygoscelis papua
## 3 Chinstrap Pygoscelis antarcticus
```

Finally, we can **combine these two data frames** with a join function.

```
penguin_bill_lengths_df %>% inner_join(latin_name_df)
```

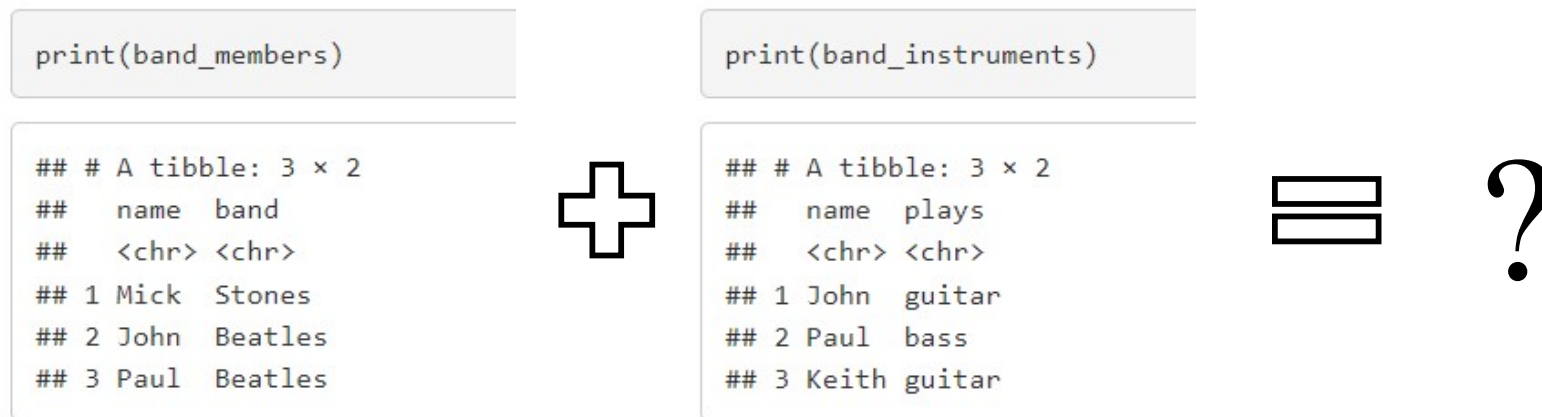
```
## # A tibble: 344 × 3
##   species    bill_length_mm latin_name
##   <fct>         <dbl> <chr>
## 1 Gentoo         59.6 Pygoscelis papua
## 2 Chinstrap      58   Pygoscelis antarcticus
## 3 Gentoo         55.9 Pygoscelis papua
## 4 Chinstrap      55.8 Pygoscelis antarcticus
## 5 Gentoo         55.1 Pygoscelis papua
## 6 Gentoo         54.3 Pygoscelis papua
## 7 Chinstrap      54.2 Pygoscelis antarcticus
## 8 Chinstrap      53.5 Pygoscelis antarcticus
## 9 Gentoo         53.4 Pygoscelis papua
## 10 Chinstrap     52.8 Pygoscelis antarcticus
## # ... with 334 more rows
```

The rows from the two data frames are merged, by matching the common columns (which is species here)

Here we have used the function *inner\_join()*, which is a type of join function. There are other types of join functions available

# Types of join functions

What happens when the set of values on the common column is not the same for both tables? For example:



**band\_members** and **band\_instruments** are two toy datasets given by the dplyr package

“Mick” only appears in “band\_members ” and “Keith” only appears in band\_instruments

Let’s rename the two data frames as x and y, respectively.

```
x = band_members  
y = band_instruments
```

There are **four basic join functions**, each of which deals with missing rows differently.

# Types of join functions 1: Inner join

An **inner join** means only rows with matching keys in both x and y are included in the result.

```
print(x)
```

```
## # A tibble: 3 × 2
##   name band
##   <chr> <chr>
## 1 Mick  Stones
## 2 John  Beatles
## 3 Paul  Beatles
```

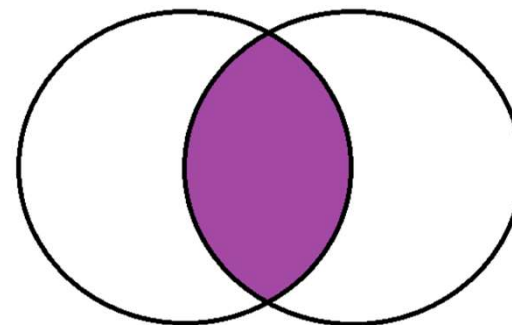
```
print(y)
```

```
## # A tibble: 3 × 2
##   name plays
##   <chr> <chr>
## 1 John  guitar
## 2 Paul  bass
## 3 Keith guitar
```

```
inner_join(x, y)
```

```
## # A tibble: 2 × 3
##   name band    plays
##   <chr> <chr>   <chr>
## 1 John  Beatles guitar
## 2 Paul  Beatles bass
```

inner join



Neither Mick nor Keith is included

# Types of join functions 2: Outer join

An **outer join** (also called a full join) means to include all rows in x with matching columns in y, then the rows of y that don't match x.

```
print(x)
```

```
## # A tibble: 3 × 2
##   name band
##   <chr> <chr>
## 1 Mick  Stones
## 2 John  Beatles
## 3 Paul  Beatles
```

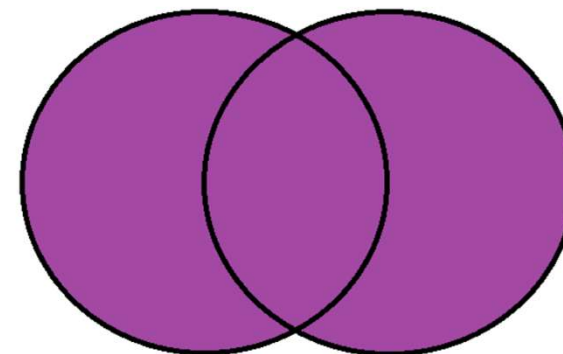
```
print(y)
```

```
## # A tibble: 3 × 2
##   name plays
##   <chr> <chr>
## 1 John  guitar
## 2 Paul  bass
## 3 Keith guitar
```

```
full_join(x, y)
```

```
## # A tibble: 4 × 3
##   name band plays
##   <chr> <chr> <chr>
## 1 Mick  Stones <NA>
## 2 John  Beatles guitar
## 3 Paul  Beatles bass
## 4 Keith <NA> guitar
```

Outer join



Both Mick and Keith are included



# Types of join functions 3: left join

A **left join** means to include all rows in x, and add matching columns from y.

```
print(x)
```

```
## # A tibble: 3 × 2
##   name band
##   <chr> <chr>
## 1 Mick  Stones
## 2 John  Beatles
## 3 Paul  Beatles
```

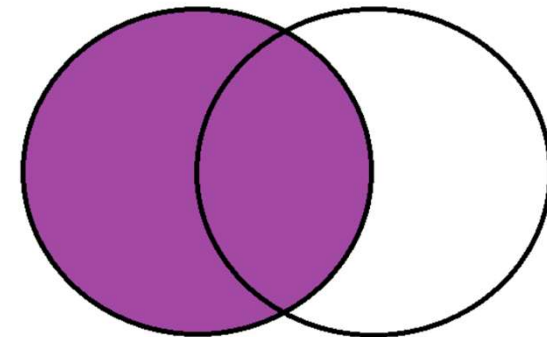
```
print(y)
```

```
## # A tibble: 3 × 2
##   name plays
##   <chr> <chr>
## 1 John  guitar
## 2 Paul  bass
## 3 Keith guitar
```

```
left_join(x, y)
```

```
## # A tibble: 3 × 3
##   name band    plays
##   <chr> <chr>   <chr>
## 1 Mick  Stones <NA>
## 2 John  Beatles guitar
## 3 Paul  Beatles bass
```

left join



Mick is included, but Keith is not



# Types of join functions 4: right join

A **right join** means to include all rows in y, and add matching columns from x.

```
print(x)
```

```
## # A tibble: 3 × 2
##   name band
##   <chr> <chr>
## 1 Mick  Stones
## 2 John  Beatles
## 3 Paul  Beatles
```


```
print(y)
```

```
## # A tibble: 3 × 2
##   name plays
##   <chr> <chr>
## 1 John  guitar
## 2 Paul  bass
## 3 Keith guitar
```

```
right_join(x, y)
```

```
## # A tibble: 3 × 3
##   name band   plays
##   <chr> <chr>   <chr>
## 1 John  Beatles guitar
## 2 Paul  Beatles bass
## 3 Keith <NA>   guitar
```

right join



Keith is included, but Mick is not

# *What we have covered today*

- We introduced basic data wrangling operations, including
  - selecting, filtering, creating and renaming columns, sorting, Summarizing, and joining multiples data frames
- We introduced the *tidyverse* and *dplyr*, basic R packages for data science
- We learned how to perform data wrangling operations using examples with R
- We explored examples with the pip operator %>%
- We explained how to use group\_by, and cross functions to obtain summary data
- We discussed different types of join functions (inner join, full join, left join, and right join)

Try the examples yourself?

The illustration, codes, and examples are included in the R Markdown file **LectureDataWrangling.Rmd** which can be downloaded via the course webpage.

Thanks for listening!

Dr Rihuan Ke

[rihuan.ke@bristol.ac.uk](mailto:rihuan.ke@bristol.ac.uk)

*Statistical Computing and Empirical Methods*  
*Unit EMATM0061, MSc Data Science*